

Lecture VII - NumPy and Pandas for Scientific Computing

Programming with Python

Dr. Tobias Vlček

Quick Recap of the last Lecture

NumPy Module

- NumPy is a fundamental package for scientific computing in Python
- It provides support for large, multi-dimensional arrays and matrices
- It also provides a wide range of mathematical functions to operate on these arrays
- Much faster than standard Python lists

Creating Arrays

- Use the `np.array(list)` function to create an array from a list
- Use the `np.zeros(shape)` function to create an array of zeros
- Use the `np.ones(shape)` function to create an array of ones
- Use the `np.arange(start, stop, step)` function to create an array of evenly spaced values
- Use the `np.linspace(start, stop, num)` function to create an array of evenly spaced values
- Use the `np.random.rand(shape)` function to create an array of random values
- Use the `np.random.randint(low, high, size)` function to create an array of random integers

Pandas Module

- Pandas is a powerful data manipulation and analysis library
- It provides data structures like DataFrames and Series for data manipulation
- It also provides tools for data cleaning, analysis, and visualization

Creating DataFrames

- Use the `pd.DataFrame(data, index, columns)` function to create a DataFrame from a dictionary
- Use the `pd.DataFrame(data, index, columns)` function to create a DataFrame from a dictionary
- Use the `pd.DataFrame(data, index, columns)` function to create a DataFrame from a dictionary

Basic Operations

- Use the `df.head()` method to display the first few rows of a DataFrame
- Use the `df.tail()` method to display the last few rows of a DataFrame
- Use the `df.info()` method to display information about a DataFrame
- Use the `df.describe()` method to display summary statistics about a DataFrame
- Use the `df.columns` attribute to access the column names of a DataFrame
- Use the `df.index` attribute to access the index of a DataFrame
- Use the `df.values` attribute to access the values of a DataFrame

Subsetting DataFrames

- Use the `df.loc[row_indexer, column_indexer]` method to access a specific element of a DataFrame
- Use the `df.iloc[row_indexer, column_indexer]` method to access a specific element of a DataFrame

Filtering DataFrames

- Use the `df[df['column'] > value]` method to filter a DataFrame
- Use the `df[df['column'].isin(values)]` method to filter a DataFrame

Grouping DataFrames

- Use the `df.groupby('column').sum()` method to group a DataFrame and calculate the sum of a column

- Use the `df.groupby('column').mean()` method to group a DataFrame and calculate the mean of a column
- Use the `df.groupby('column').count()` method to group a DataFrame and count the number of elements in a column
- Use the `df.groupby('column').size()` method to group a DataFrame and count the number of elements in a column

Excel Files

- Excel files can be read using the `pd.read_excel(file_path)` function
- Excel files can be written using the `df.to_excel(file_path)` method