Lecture VII - Pandas and AI

Programming with Python

Dr. Tobias Vlćek

Kühne Logistics University Hamburg - Fall 2025

Quick Recap of the last Lecture

What is NumPy?

- NumPy is a package for scientific computing in Python
- Provides multi-dimensional arrays and matrices
- Much faster than Python lists for numerical operations
- Operations are implemented in C and C++

. . .

> 💡 Tip
>
> NumPy arrays are stored in contiguous memory blocks, making operations very efficient.

Creating Arrays

- Core data structure is the `ndarray`
- Can create arrays from lists, tuples, or other data structures
- Special functions like:
  - ‣ `np.zeros()` for arrays of zeros
  - ‣ `np.random.rand()` for random values
  - ‣ `np.arange()` for evenly spaced values
  - ‣ `np.linspace()` for linearly spaced values

Working with Arrays

- Support for multi-dimensional operations
- Common operations:
  - ‣ Element-wise arithmetic (`+`, `-`, `*`, `/`)
  - ‣ Array indexing and slicing
  - ‣ Shape manipulation (`reshape`, `flatten`)
  - ‣ Sorting and transposing

. . .

> 💡 Tip
>
> NumPy operations are vectorized, meaning they operate on entire arrays at once rather than element by element.

NumPy in Action I

Task: Complete the following task:

```python
# TODO: Create an array with 10 evenly spaced numbers over the interval
from 0 to 73.

import numpy as np
# YOUR CODE HERE
```

. . .

> ℹ Note
>
> Note, that you can always use the `help()` function to get more information about a function. But be sure to import the package first, otherwise you will get an error. To quit the help page, press `q`.

NumPy in Action II

Task: Complete the following task:

```python
# TODO: Take the following 3x3 array and reduce it to a 1D array.

import numpy as np
array = np.array([[1, 2, 3], [4, 5, 6], [7, 8, 9]])
# YOUR CODE HERE
```

Pandas Basics

What is Pandas?

- Pandas is a data manipulation and analysis library
- It provides data structures like DataFrames and Series
- Tools for data cleaning, analysis, and visualization
- It can also be used to work with Excel files!

How to install Pandas

- In the last lecture, we have installed it with `uv install pandas`
- Now, import the package `import pandas as pd`

. . .

> **i Note**
>
> You can also use a different abbreviation, but `pd` is the most common one.

Creating DataFrames

- DataFrames behave quite similar to Numpy arrays
- But they have row and column labels

...

```python
import pandas as pd
df = pd.DataFrame({ # DataFrame is created from a dictionary
    "Name": ["Tobias", "Robin", "Nils", "Nikolai"],
    "Kids": [2, 1, 0, 0],
    "City": ["Oststeinbek", "Oststeinbek", "Hamburg", "Lübeck"],
    "Salary": [3000, 3200, 4000, 2500]}); print(df)
```

```
      Name  Kids         City  Salary
0   Tobias     2  Oststeinbek    3000
1    Robin     1  Oststeinbek    3200
2     Nils     0      Hamburg    4000
3  Nikolai     0       Lübeck    2500
```

Reading from CSV Files

```python
df = pd.read_csv("supplementary/lec_08/employees.csv") # Reads the CSV file
print(df)
```

```
       Name  Age Department        Position  Salary
0     Alice   30         HR         Manager   50000
1       Bob   25         IT       Developer   60000
2   Charlie   28    Finance         Analyst   55000
3     David   35  Marketing       Executive   52000
4       Eve   32      Sales  Representative   48000
5     Frank   29         IT       Developer   61000
6     Grace   31         HR       Assistant   45000
7      Hank   27    Finance         Analyst   53000
8       Ivy   33  Marketing         Manager   58000
9      Jack   26      Sales  Representative   47000
10     Kara   34         IT       Developer   62000
11      Leo   30         HR         Manager   51000
12     Mona   28    Finance         Analyst   54000
13     Nina   35  Marketing       Executive   53000
14    Oscar   32      Sales  Representative   49000
15     Paul   29         IT       Developer   63000
16    Quinn   31         HR       Assistant   46000
17     Rita   27    Finance         Analyst   52000
18      Sam   33  Marketing         Manager   59000
19     Tina   26      Sales  Representative   48000
```

```
20      Uma   34        IT     Developer   64000
21    Vince   30        HR       Manager   52000
22     Walt   28   Finance       Analyst   55000
23     Xena   35  Marketing    Executive   54000
24     Yara   32     Sales  Representative  50000
25     Zane   29        IT     Developer   65000
26     Anna   31        HR     Assistant   47000
27      Ben   27   Finance       Analyst   53000
28    Cathy   33  Marketing      Manager   60000
29    Dylan   26     Sales  Representative  49000
30     Ella   34        IT     Developer   66000
31     Finn   30        HR       Manager   53000
32     Gina   28   Finance       Analyst   56000
33     Hugo   35  Marketing    Executive   55000
34     Iris   32     Sales  Representative  51000
35     Jake   29        IT     Developer   67000
36     Kyla   31        HR     Assistant   48000
37     Liam   27   Finance       Analyst   54000
38      Mia   33  Marketing      Manager   61000
39     Noah   26     Sales  Representative  50000
40    Olive   34        IT     Developer   68000
41     Pete   30        HR       Manager   54000
42   Quincy   28   Finance       Analyst   57000
43     Rose   35  Marketing    Executive   56000
44    Steve   32     Sales  Representative  52000
45     Tara   29        IT     Developer   69000
46     Umar   31        HR     Assistant   49000
47     Vera   27   Finance       Analyst   55000
48     Will   33  Marketing      Manager   62000
49     Zara   26     Sales  Representative  51000
```

Basic Operations

- Use the `df.head()` method to display the first 5 rows
- Use the `df.tail()` method to display the last 5 rows

...

```python
df = pd.read_csv("supplementary/lec_08/employees.csv")
print(df.tail())
```

```
    Name  Age Department       Position  Salary
45  Tara   29         IT      Developer   69000
46  Umar   31         HR      Assistant   49000
47  Vera   27    Finance        Analyst   55000
48  Will   33  Marketing        Manager   62000
49  Zara   26      Sales  Representative  51000
```

Information about the DataFrame

- Use `df.info()` to display information about a DataFrame

...

```
df = pd.read_csv("supplementary/lec_08/employees.csv")
print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50 entries, 0 to 49
Data columns (total 5 columns):
 #   Column      Non-Null Count  Dtype
---  ------      --------------  -----
 0   Name        50 non-null     object
 1   Age         50 non-null     int64
 2   Department  50 non-null     object
 3   Position    50 non-null     object
 4   Salary      50 non-null     int64
dtypes: int64(2), object(3)
memory usage: 2.1+ KB
None
```

Statistics about a DataFrame

- Use `df.describe()` to display summary statistics
- Use the `df.index` attribute to access the index

…

```
df = pd.read_csv("supplementary/lec_08/employees.csv")
print(df.describe())
```

```
             Age        Salary
count  50.000000     50.000000
mean   30.320000  54980.000000
std     2.958488   6175.957333
min    25.000000  45000.000000
25%    28.000000  50250.000000
50%    30.000000  54000.000000
75%    33.000000  59750.000000
max    35.000000  69000.000000
```

Filtering DataFrames

- Use `df['column_name']` to access a column
- Use the `df[df['column'] > value]` method to filter

…

```
df = pd.read_csv("supplementary/lec_08/employees.csv")
df_high_salary = df[df['Salary'] >= 67000]
print(df_high_salary)
print(df_high_salary.iloc[2]["Name"]) #Access the third row and the "Name"
column
print(df_high_salary.loc[40]["Name"]) #Access the label 40 and the "Name"
column
```

```
      Name  Age Department   Position  Salary
35    Jake   29         IT  Developer   67000
40   Olive   34         IT  Developer   68000
45    Tara   29         IT  Developer   69000
Tara
Olive
```

Filtering in Action

Task: Complete the following task:

```
# TODO: Load the employees.csv located in the git repository into a
DataFrame
# First, filter the DataFrame for employees with a manager position
# Then, print the average salary of the remaining employees
# Finally, print the name of the employee with the lowest salary
```

...

> **i** Note
>
> Note, that we can use the `mean()` method on the `Salary` column, as it is a numeric column. In addition, we can use the `min()` method on the `Salary` column to find the lowest salary.

Grouping DataFrames

Grouping

- Grouping is a powerful feature of Pandas
- Groups data by one or more columns
- And then perform operations
- Syntax is `df.groupby('column').method()`

...

```
df = pd.read_csv("supplementary/lec_08/employees.csv")
df.groupby(['Position']).sum() # Sum per position
```

Position	Name	Age	Department	Salary
Analyst	CharlieHankMonaR-itaWaltBenGina-LiamQuincyVera	275	FinanceFinanceFi-nanceFinanceFinance-FinanceFina...	544000
Assistant	GraceQuinnAnnaKy-laUmar	155	HRHRHRHRHR	235000

	Position	Name	Age	Department	Salary
Developer		BobFrankKara-PaulUmaZaneElla-JakeOliveTara	306	ITITITITITITITITIT	645000
Executive		DavidNinaXe-naHugoRose	175	MarketingMarketing-MarketingMarketing-Marketing	270000
Manager		AliceIvyLeoSamVince-CathyFinnMiaPeteWill	315	HRMarketingHRMar-ketingHRMarket-ingHRMarketingHR...	560000
Representative		EveJackOscarTina-YaraDylanIrisNoah-SteveZara	290	SalesSalesSalesSa-lesSalesSalesSalesSa-lesSalesS...	495000

Grouping Numeric Columns
- To prevent errors, we can select numeric columns first
- Afterwards, perform the operation on the selected columns
- Helps to avoid errors when grouping by non-numeric columns
- Or drop columns by `df.drop(columns=["column"])`

...

```python
df = pd.read_csv("supplementary/lec_08/employees.csv")
numeric_cols = df.select_dtypes(include=['number']).columns
print(df.groupby("Position")[numeric_cols].sum())
```

```
                Age  Salary
Position
Analyst         275  544000
Assistant       155  235000
Developer       306  645000
Executive       175  270000
Manager         315  560000
Representative  290  495000
```

Grouping by Multiple Columns
- Group by multiple columns `['column1', 'column2']`
- You can use lists or tuples to specify multiple columns

...

```python
df = pd.read_csv("supplementary/lec_08/employees.csv")
df = df.drop(columns=["Name"])
```

```
# Max per position and department
df.groupby(['Position', "Department"]).max()
```

		Age	Salary
Position	Department		
Analyst	Finance	28	57000
Assistant	HR	31	49000
Developer	IT	34	69000
Executive	Marketing	35	56000
Manager	HR	30	54000
	Marketing	33	62000
Representative	Sales	32	52000

Grouping with Aggregations

- We can use different aggregation functions:
  - ‣ `sum()`: sum of the values
  - ‣ `mean()`: mean of the values
  - ‣ `max()`: maximum of the values
  - ‣ `min()`: minimum of the values
  - ‣ `count()`: count of the values

Pandas in Action

Task: Complete the following task:

```
# TODO: Load the employees.csv again into a DataFrame
# First, group by the "Position" column and count the employees per
position
# Then, group by the "Department" column and calculate the mean of all
other columns per department
df = pd.read_csv("supplementary/lec_08/employees.csv")
# Your code here
```

Combining DataFrames

Concatenating DataFrames

- `pd.concat()` to concatenate along shared columns

```
df1 = pd.DataFrame({"A": [1, 2, 3], "B": [4, 5, 6]})
df2 = pd.DataFrame({"A": [7, 8, 9], "B": [10, 11, 12]})
df = pd.concat([df1, df2])
print(df)
```

```
     A   B
0    1    4
1    2    5
2    3    6
0    7   10
1    8   11
2    9   12
```

Joining DataFrames

- Use `pd.join()` to join DataFrames along columns
- Joining is done on the index by default!

```
df1 = pd.DataFrame({"A": [1, 2, 3], "B": [4, 5, 6]}, index=['x', 'y', 'z'])
df2 = pd.DataFrame({"C": [7, 8, 9], "D": [10, 11, 12]}, index=['z', 'y',
'w'])
df = df1.join(df2)
print(df)
```

```
     A  B    C     D
x    1  4  NaN   NaN
y    2  5  8.0  11.0
z    3  6  7.0  10.0
```

Merging DataFrames on Columns

- `pd.merge(df_name, on='column', how='type')`
- merge DataFrames along shared columns
- `how` specifies the type of merge
    - `inner`: rows with matching keys in both DataFrames
    - `outer`: rows from both are kept, missing values are filled
    - `left`: rows from the left are kept, missing values are filled
    - `right`: rows from right are kept, missing values are filled

Outer Merge

```
df3 = pd.DataFrame({"A": [1, 2, 3], "B": [4, 5, 6]})
df4 = pd.DataFrame({"A": [2, 3, 4], "C": [7, 8, 9]})
df_merged = df3.merge(df4, on="A", how="outer")
print(df_merged)
```

```
     A    B    C
0    1  4.0  NaN
1    2  5.0  7.0
2    3  6.0  8.0
3    4  NaN  9.0
```

Merging in Action

Task: Complete the following task:

```python
df1 = pd.DataFrame({
    "Name": ["John", "Alice", "Bob", "Carol"],
    "Department": ["Sales", "IT", "HR", "Sales"],
    "Salary": [50000, 60000, 55000, 52000]})
df2 = pd.DataFrame({
    "Name": ["Alice", "Bob", "Dave", "Eve"],
    "Position": ["Developer", "Manager", "Analyst", "Developer"],
    "Years": [5, 8, 3, 4]})

# TODO: Merge the two DataFrames on the "Name" column
# Try different types of merges (inner, outer, left, right)
# Observe and describe the differences in the results
```

Working with Excel Files

Reading Excel Files

- Read using the `pd.read_excel(file_path)` function
- Write using the `df.to_excel(file_path)` method

...

```python
import pandas as pd
df = pd.read_csv("supplementary/lec_08/employees.csv")
df.to_excel("supplementary/lec_08/employees.xlsx", index=False)
```

...

> **i Note**
>
> Note, that you likely need to install the `openpyxl` package to be able to write Excel files, as it handles the file format.

Advanced Excel file handling

We can also specify the sheet name when reading and writing

```python
# Writes to the Employees sheet and does not include row indices
df.to_excel("supplementary/lec_08/employees.xlsx", sheet_name="Employees",
index=False)
```

...

```python
# Reads from the Employees sheet
df = pd.read_excel("supplementary/lec_08/employees.xlsx",
sheet_name="Employees")
print(df.head())
```

```
        Name  Age Department        Position  Salary
0      Alice   30         HR         Manager   50000
1        Bob   25         IT       Developer   60000
2    Charlie   28    Finance         Analyst   55000
3      David   35  Marketing       Executive   52000
4        Eve   32      Sales  Representative   48000
```

Excel in Action

Task: Complete the following task:

```
# TODO: Load the temperatures.xlsx file into a DataFrame
# Look at the first few rows of the DataFrame
# Then, print the average temperature per city
```

Melting DataFrames

Melting

- Sometimes, you want to transform a DataFrame
- Instead of wide format, you want long format
- This is useful for certain types of visualizations
- And when working with time series data

...

Question: Anybody ever heard of the terms?

Wide Format

For example, the following DataFrame is in wide format:

```
          Date  Hamburg  Los_Angeles  Tokyo
0   2024-03-01     12.0         18.2   14.8
1   2024-03-02      9.8         23.0   17.6
2   2024-03-03      7.6         20.3   16.0
3   2024-03-04     10.1         21.1   13.4
4   2024-03-05     11.2         18.5   15.1
..         ...      ...          ...    ...
87  2024-05-27     12.4         24.5   24.9
88  2024-05-28     17.8         20.6   22.3
89  2024-05-29     16.2         20.4   20.2
90  2024-05-30     15.5         20.7   21.7
91  2024-05-31     12.6         22.0   22.9

[92 rows x 4 columns]
```

Long Format

The melting process transforms it into the following long format:

```
          Date     City  Temperature
0   2024-03-01  Hamburg         12.0
1   2024-03-02  Hamburg          9.8
2   2024-03-03  Hamburg          7.6
3   2024-03-04  Hamburg         10.1
4   2024-03-05  Hamburg         11.2
..         ...      ...          ...
271 2024-05-27    Tokyo         24.9
272 2024-05-28    Tokyo         22.3
273 2024-05-29    Tokyo         20.2
274 2024-05-30    Tokyo         21.7
275 2024-05-31    Tokyo         22.9

[276 rows x 3 columns]
```

How to melt DataFrames

- Use `pd.melt()` to transform from wide to long
- Parameters:
  - ‣ `id_vars`: columns to keep
  - ‣ `var_name`: name of the new column that will contain the names of the original columns
  - ‣ `value_name`: name of the new column that will contain the values of the original columns

...

```python
df = pd.read_csv("supplementary/lec_08/employees.csv")
df = pd.melt(df, id_vars=['Position'], var_name='Variables',
value_name='Values')
print(df)
```

```
           Position Variables     Values
0           Manager      Name      Alice
1         Developer      Name        Bob
2           Analyst      Name    Charlie
3         Executive      Name      David
4    Representative      Name        Eve
..              ...       ...        ...
195       Developer    Salary      69000
196       Assistant    Salary      49000
197         Analyst    Salary      55000
198         Manager    Salary      62000
199  Representative    Salary      51000

[200 rows x 3 columns]
```

Melting in Action

Task: Complete the following task:

```
# TODO: Load and transform the temperatures.xlsx file by melting it
# Expected output format:
#          Date         City  Temperature
# 0  2024-03-01     Hamburg          7.2
# 1  2024-03-01 Los_Angeles         18.5
# 2  2024-03-01       Tokyo         12.3
# Then, print the maximum temperature per city by grouping by the "City"
column
```

Programming with AI

Using AI to generate code
- Coding by hand is not the only way to generate code
- Most likely, a lot of you have already used ChatGPT

...

How do

Large Language

Models work?

Photo by Taylor Vick on Unsplash

Large Language Models (LLMs)
- Think of them like advanced pattern recognition systems
- They have "read" massive amounts of text
- Books, websites, articles, code, and more
- Text is broken into tokens, parts of words or punctuation
- Based on patterns, they can generate new text

Training LLMs
- Imagine learning a language by reading millions of books
- Learns patterns in how words and ideas connect via tokens
- Interconnected nodes with weights representing patterns
- During training, these weights are adjusted
- Once trained, applying them takes much less ressources

Pattern Recognition
- Not like a search engine!
- When asked, it looks for relevant patterns it learned
- Like having a huge library in its "memory" to draw from
- It can find patterns between concepts and your question
- Knows only limited text at once (context window)

Probability based responses
- After each written token, it predicts "what should come next?"

- Like a advanced version of the word prediction on your phone
- Chooses the most likely next token based on training
- But can't actually "think" or "understand" like humans

Limitations
- No true understanding of cause and effect
- Sometimes makes mistakes or "hallucinates"
- Mostly only knows what it was trained on
- Can reflect biases present in training data
- No emotional understanding (but can simulate responses!)

Impact on Jobs
- Question: What do you think about their impact on jobs?
- Question: What are the implications for us?
- Question: Can we use them to our advantage?

. . .

> ⚠️ Warning
>
> If you use free models, be aware that your prompts are going to be used by the providers and are not private. But for learning and experimenting, this should be no issue.

AI Coding Partner

(Current) Choices for Programmers
- Github Copilot: Integrated into VS Code by Microsoft
- Cursor: Fork of VS Code with AI assistance built in
- Aider: Chat interface for AI to write code in the terminal
- Zed: Lightweight IDE with AI features
- Claude Code: Coding-Assistant in the terminal

. . .

> 💡 Tip
>
> Currently, Zed in combination with Claude Code is my favorite one. But this might change in the future, as there is a lot of competition in this space.

GitHub Copilot
GitHub Copilot helps you write code faster and with less effort.

. . .

Think of it as:

- An autocomplete for entire lines or blocks of code
- A coding assistant that understands context
- A learning tool that shows you coding patterns

. . .

> **ⓘ Note**
>
> Copilot uses AI trained on billions of lines of public code to suggest completions based on what you're typing.

Get Free Access

GitHub Student Developer Pack gives you free Copilot access!

. . .

How to get it:

1. Go to education.github.com/pack
2. Sign up with your university email
3. Verify your student status
4. Wait for approval (usually 1-2 days)
5. Login into your account in VS Code

. . .

> **ⓘ Note**
>
> You'll need a GitHub account. Create one at github.com if you don't have one.

Asking for help

Task: Paste the following prompt in to the chat:

Can you please write me a small number guessing game in python? It should work for one player in the terminal. The player should guess a number between 1-10 and get hints about whether his guess was too large or too small. After 3 tries, end the game if he didn't succeed with a nice message.

. . .

Copy the generated code and paste it into your file.

More on Copilot

- While working with Copilot, it will suggest you code changes
- You can accept or reject them
- The rest you will learn by doing!

. . .

> **ℹ Note**
>
> And that's it for todays lecture!
> You now have the basic knowledge to start working with tabular data and AI!.

Literature

Interesting Books

- Downey, A. B. (2024). Think Python: How to think like a computer scientist (Third edition). O'Reilly. Link to free online version
- Elter, S. (2021). Schrödinger programmiert Python: Das etwas andere Fachbuch (1. Auflage). Rheinwerk Verlag.

. . .

For more interesting literature to learn more about Python, take a look at the literature list of this course.