

Lecture XI - Project Work I

Programming with Python

Dr. Tobias Vlček

Kühne Logistics University Hamburg - Fall 2026

Episode 11: The Build

The shop has a new owner

MunchCorp runs the bakery now. Tobi kept his job, the investor kept her spreadsheets, and her next due-diligence visit already has a date: **Session XIII**, when your projects run in this room.

...

Until then, the shop rule stands: pull before you start, push when you stop, and let the commit history do the talking.

The Mark

Could anyone tell?

- Your README carries an AI-disclosure section. Suppose you skipped it: could anyone prove that an AI wrote part of your work?
- Since August 2026, EU rules require providers to mark AI content machine-readably, and the major ones do
- The trick behind it is simple enough for three slides

How a watermark hides in plain text

- Before each word, the model splits its vocabulary into a **green** and a **red** half, reshuffled at every step
- It then softly prefers green words. The text still reads completely normally
- By chance, a human writes about half green words. A model writes far more

You already know the mechanism

- Session VI: the model picks a likely next token, never *the* next token
- A watermark is that choice, steered on purpose: green wins a little more often than chance
- Same next-token choice, one hidden thumb on the scale

Who can actually detect it

- Detection is counting: one paragraph of marked text can push chance below 1 in 10 trillion
- But you need the secret key that decided which half was green

- In practice, only the provider can run the test

What a mark proves (and does not)

- Proves: a model processed this text
- Does not prove: who did the work, or whether a human verified it
- Heavy editing can wash a mark out, and short snippets may not carry one, so an absence proves nothing

...

Presence without disclosure, though, looks exactly like what it is.

Disclosure is part of the work

- The rules converged everywhere: AI cannot be an author, and its use gets disclosed
- Your README section says which tool, for what, and what you verified yourself
- That is standard practice, not a confession

The Work Menu

Today

- **Pull before you start, push when you stop.** The remote is your backup; commit as you go
- Progress today means one working increment on `main`, with commits that tell the story
- I am in the room: use me for blockers, design questions, and “is this scope realistic?”
- Use the agent where it helps, and note down what it did for your disclosure section

When you're stuck

Tip

A red traceback you have never seen? Read the last line first, exactly as in Part I. Then, and only then, the agent.

Wrap-up

Three things to remember

- AI text carries a watermark now: fluent, normal-looking, detectable by the provider
- Disclosure is part of the work: which tool, what for, what you verified
- Small commits, pushed often, are the cheapest insurance either of you can buy

...

Next week: your agent keeps making the same mistakes. Time to **teach it once** instead of correcting it every session.